diffst - A tool for Data Integration Testing

Francois Marais

francois@busii.com

Outline

- Data Integration Applications
- Functional Testing Aproaches
- Solution provided by Market Leader
- Demo
- diffst - Features and limitations
- Questions and Discussion

Data Integration/Migration Applications

- Once-off projects
  - Data Migration - hardware, software, database upgrade/relocation
  - Application Migration (CRM/ERP)
  - Take-on of new data sets (eg corporate mergers/acquisitions)
- Systems and Applications
  - Data Warehousing
  - Business Intelligence, Big Data,...
  - Customer Loyalty Programme Engines

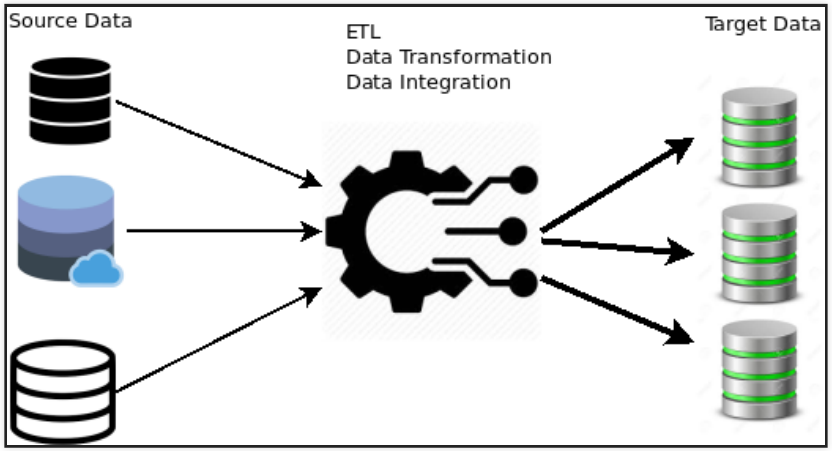

Data Integration Applications - Typical Risks

- Data not extracted from Source as expected
  - Wrong time period
  - Not all expected data extracted
  - Extracted at wrong time
- Expected business rules not applied
  - Fields truncated
  - Data Types misread
- Target data Incomplete/Duplicated/Inconsistent..
- Wrong grain

Coverage Challenges

- Large actual data volumes and
- Business logic complexity - per column
- Implies many actual data points available for testing
- ... and they need to be tested - more important than synthetic test data

Common functional testing checks

- SELECT data from Source
- Apply business rules
- SELECT expected data from Target
- Compare and evaluate
- ⇒ Automation opportunity

Sampling: Comparing individual data items

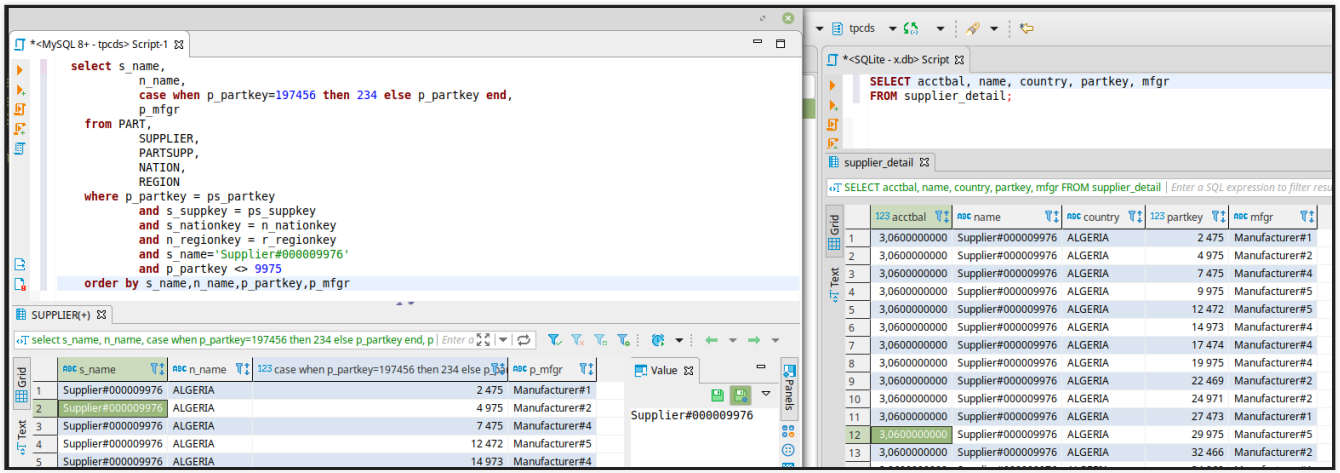

- Useful for testing specific data items - Exploratory value
- But not effective in testing large data volumes
- How to build into regression, test suites?

Aggregated SQL

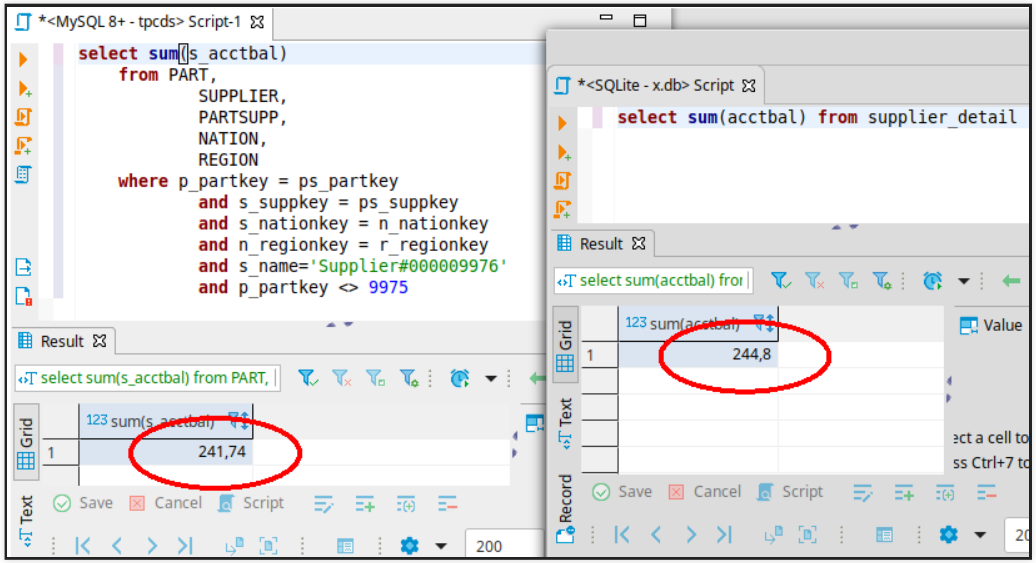

- Can only test few columns at a time
- Slow to investigate source-target discrepancies

MINUS Queries

Table_1 Name: Yankees1

f_name	l_name	Position
Babe	Ruth	LF
Louis	Gehrig	1B
Joe	DiMaggio	CF
Mickey	Mantle	CF
Derek	Jeter	SS
Yogi	Berra	C
Willie	Randolph	2B
Graig	Nettles	3B
CC	Sabathia	P

Table_2 Name: Yankees2

firstName	lastName	Positions
Babe	Ruth	LF
Louis	Gehrig	1B
Joe	DiMaggio	CF
Mickey	Mantle	CF
Derek	Jeter	SS
Yogi	Berra	C
Willie	Randolph	2B
Andy	Pettitte	P

So the first minus query for these tables is Table_1 MINUS Table_2:

(SELECT f_name, l_name, position FROM Yankees1
MINUS
SELECT firstName, lastName, Positions FROM Yankees2)

The result set should be the rows we have highlighted in RED below:

Graig	Nettles	3B
CC	Sabathia	P

Then you need to subtract Table_2 MINUS Table_1:

(SELECT firstName, lastName, Positions FROM Yankees2
MINUS
SELECT f_name, l_name, position FROM Yankees1)

The result set should be the rows we have highlighted in GREEN below:

Andy	Pettitte	P

- Slow - have to run a query on each side separately - then compare

MINUS Queries

Table_1 Name: Yankees1

f_name	l_name	Position
Babe	Ruth	LF
Louis	Gehrig	1B
Joe	DiMaggio	CF
Mickey	Mantle	CF
Derek	Jeter	SS
Yogi	Berra	C
Willie	Randolph	2B
Graig	Nettles	3B
CC	Sabathia	P

Table_2 Name: Yankees2

firstName	lastName	Positions
Babe	Ruth	LF
Louis	Gehrig	1B
Joe	DiMaggio	CF
Mickey	Mantle	CF
Derek	Jeter	SS
Yogi	Berra	C
Willie	Randolph	2B
Andy	Pettitte	P

So the first minus query for these tables is Table_1 MINUS Table_2:

(SELECT f_name, l_name, position FROM Yankees1
MINUS
SELECT firstName, lastName, Positions FROM Yankees2)

The result set should be the rows we have highlighted in RED below:

Graig	Nettles	3B
CC	Sabathia	P

Then you need to subtract Table_2 MINUS Table_1:

(SELECT firstName, lastName, Positions FROM Yankees2
MINUS
SELECT f_name, l_name, position FROM Yankees1)

The result set should be the rows we have highlighted in GREEN below:

Andy	Pettitte	P

- Source and Target must be on same database
- Slow - have to run a query on each side separately - then compare

Source

	Col1	Col2	Col3
Row1	xxxxx	xxxxx	xxxxx
Row2	xxxxx	xxxxx	xxxxx
Row3	xxxxx	xxxxx	xxxxx
Row4	xxxxx	xxxxx	xxxxx
Row5	xxxxx	xxxxx	xxxxx
Row6	xxxxx	xxxxx	xxxxx
Row7	xxxxx	xxxxx	xxxxx
Row8	xxxxx	xxxxx	xxxxx
Row9	xxxxx	xxxxx	xxxxx
Row10	xxxxx	xxxxx	xxxxx
Row11	xxxxx	xxxxx	xxxxx
Row12	xxxxx	xxxxx	xxxxx
Row13	xxxxx	xxxxx	xxxxx
....	xxxxx	xxxxx	xxxxx

Target

	Col1	Col2	Col3
Row1	xxxxx	xxxxx	xxxxx
Row2	xxxxx	xxxxx	xxxxx
Row3	xxxxx	xxxxx	xxxxx
Row4	xxxxx	xxxxx	xxxxx
Row5	xxxxx	xxxxx	xxxxx
Row6	xxxxx	xxxxx	xxxxx
Row7	xxxxx	xxxxx	xxxxx
Row8	xxxxx	xxxxx	xxxxx
Row9	xxxxx	xxxxx	xxxxx
Row10	xxxxx	xxxxx	xxxxx
Row11	xxxxx	xxxxx	xxxxx
Row12	xxxxx	xxxxx	xxxxx
Row13	xxxxx	xxxxx	xxxxx
....	xxxxx	xxxxx	xxxxx

Manual Detailed Comparison

Source

	Col1	Col2	Col3
Row1	xxxxx	xxxxx	xxxxx
Row2	xxxxx	xxxxx	xxxxx
Row3	xxxxx	xxxxx	xxxxx
Row4	xxxxx	xxxxx	xxxxx
Row5	xxxxx	xxxxx	xxxxx
Row6	xxxxx	xxxxx	xxxxx
Row7	xxxxx	xxxxx	xxxxx
Row8	xxxxx	xxxxx	xxxxx
Row9	xxxxx	xxxxx	xxxxx
Row10	xxxxx	xxxxx	xxxxx
Row11	xxxxx	xxxxx	xxxxx
Row12	xxxxx	xxxxx	xxxxx
Row13	xxxxx	xxxxx	xxxxx
....	xxxxx	xxxxx	xxxxx

Target

	Col1	Col2	Col3
Row1	xxxxx	xxxxx	xxxxx
Row2	xxxxx	xxxxx	xxxxx
Row3	xxxxx	xxxxx	xxxxx
Row4	xxxxx	xxxxx	xxxxx
Row5	xxxxx	xxxxx	xxxxx
Row6	xxxxx	xxxxx	xxxxx
Row7	xxxxx	xxxxx	xxxxx
Row8	xxxxx	xxxxx	xxxxx
Row9	xxxxx	xxxxx	xxxxx
Row10	xxxxx	xxxxx	xxxxx
Row11	xxxxx	xxxxx	xxxxx
Row12	xxxxx	xxxxx	xxxxx
Row13	xxxxx	xxxxx	xxxxx
....	xxxxx	xxxxx	xxxxx

Group By Query Comparison

Source

	Col1	Col2	Col3
Row1	xxxxx	xxxxx	xxxxx
Row2	xxxxx	xxxxx	xxxxx
Row3	xxxxx	xxxxx	xxxxx
Row4	xxxxx	xxxxx	xxxxx
Row5	xxxxx	xxxxx	xxxxx
Row6	xxxxx	xxxxx	xxxxx
Row7	xxxxx	xxxxx	xxxxx
Row8	xxxxx	xxxxx	xxxxx
Row9	xxxxx	xxxxx	xxxxx
Row10	xxxxx	xxxxx	xxxxx
Row11	xxxxx	xxxxx	xxxxx
Row12	xxxxx	xxxxx	xxxxx
Row13	xxxxx	xxxxx	xxxxx
....	xxxxx	xxxxx	xxxxx

Target

	Col1	Col2	Col3
Row1	xxxxx	xxxxx	xxxxx
Row2	xxxxx	xxxxx	xxxxx
Row3	xxxxx	xxxxx	xxxxx
Row4	xxxxx	xxxxx	xxxxx
Row5	xxxxx	xxxxx	xxxxx
Row6	xxxxx	xxxxx	xxxxx
Row7	xxxxx	xxxxx	xxxxx
Row8	xxxxx	xxxxx	xxxxx
Row9	xxxxx	xxxxx	xxxxx
Row10	xxxxx	xxxxx	xxxxx
Row11	xxxxx	xxxxx	xxxxx
Row12	xxxxx	xxxxx	xxxxx
Row13	xxxxx	xxxxx	xxxxx
....	xxxxx	xxxxx	xxxxx

Test Tool – Detailed Comparison of entire data sets

Source

	Col1	Col2	Col3
Row1	xxxxx	xxxxx	xxxxx
Row2	xxxxx	xxxxx	xxxxx
Row3	xxxxx	xxxxx	xxxxx
Row4	xxxxx	xxxxx	xxxxx
Row5	xxxxx	xxxxx	xxxxx
Row6	xxxxx	xxxxx	xxxxx
Row7	xxxxx	xxxxx	xxxxx
Row8	xxxxx	xxxxx	xxxxx
Row9	xxxxx	xxxxx	xxxxx
Row10	xxxxx	xxxxx	xxxxx
Row11	xxxxx	xxxxx	xxxxx
Row12	xxxxx	xxxxx	xxxxx
Row13	xxxxx	xxxxx	xxxxx
....	xxxxx	xxxxx	xxxxx

Target

	Col1	Col2	Col3
Row1	xxxxx	xxxxx	xxxxx
Row2	xxxxx	xxxxx	xxxxx
Row3	xxxxx	xxxxx	xxxxx
Row4	xxxxx	xxxxx	xxxxx
Row5	xxxxx	xxxxx	xxxxx
Row6	xxxxx	xxxxx	xxxxx
Row7	xxxxx	xxxxx	xxxxx
Row8	xxxxx	xxxxx	xxxxx
Row9	xxxxx	xxxxx	xxxxx
Row10	xxxxx	xxxxx	xxxxx
Row11	xxxxx	xxxxx	xxxxx
Row12	xxxxx	xxxxx	xxxxx
Row13	xxxxx	xxxxx	xxxxx
....	xxxxx	xxxxx	xxxxx

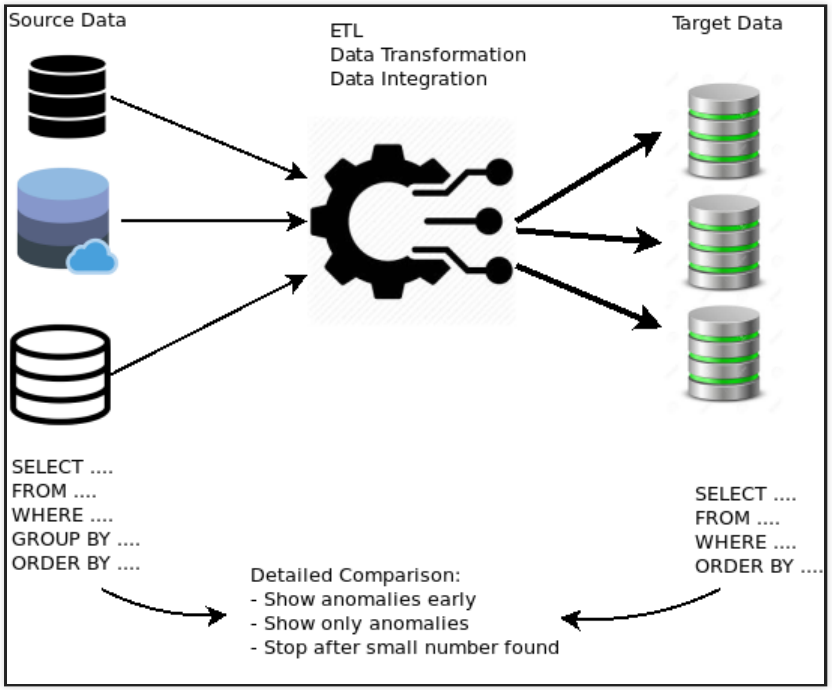

Requirement for Test Tool

The Market Leader: QuerySurge

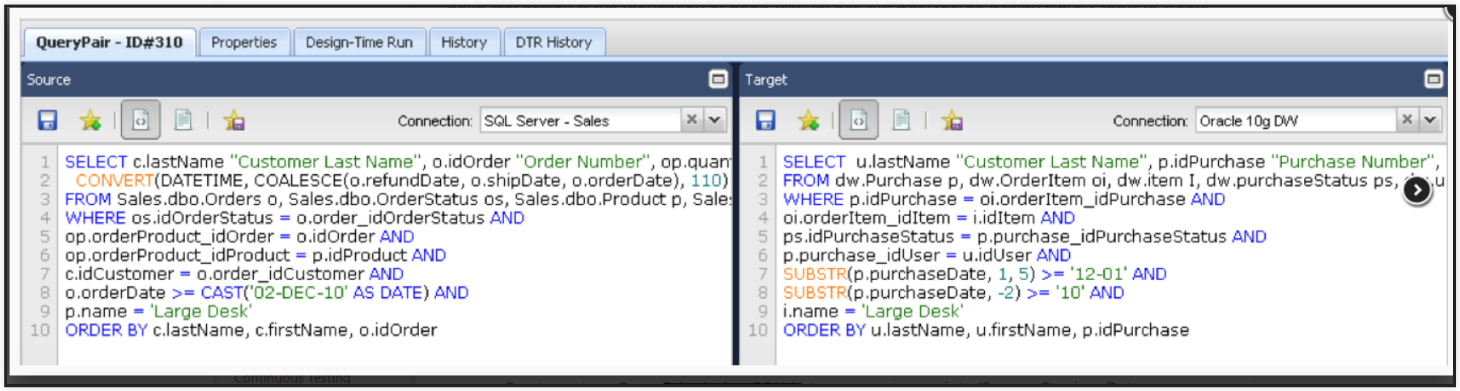

QuerySurge - Results

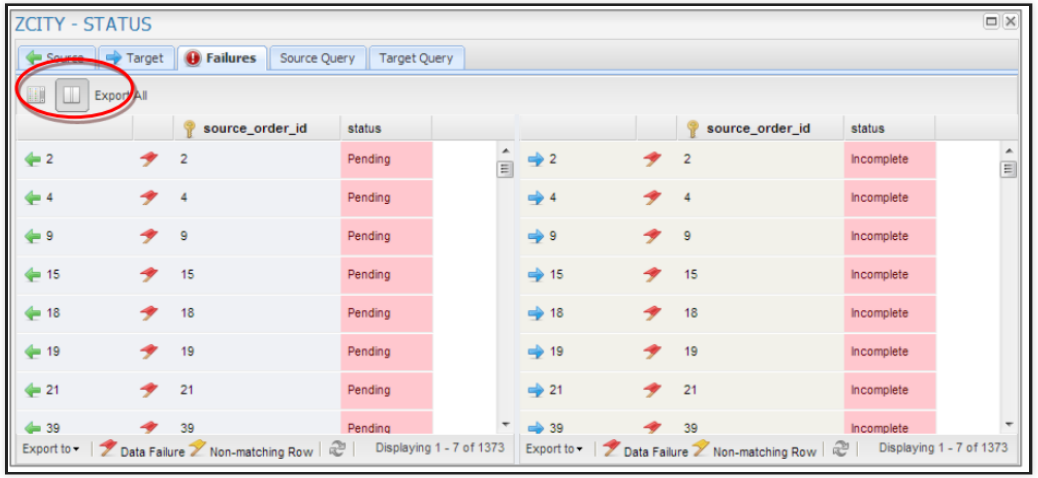

It is All-Singing, All-Dancing

Has a Scheduler

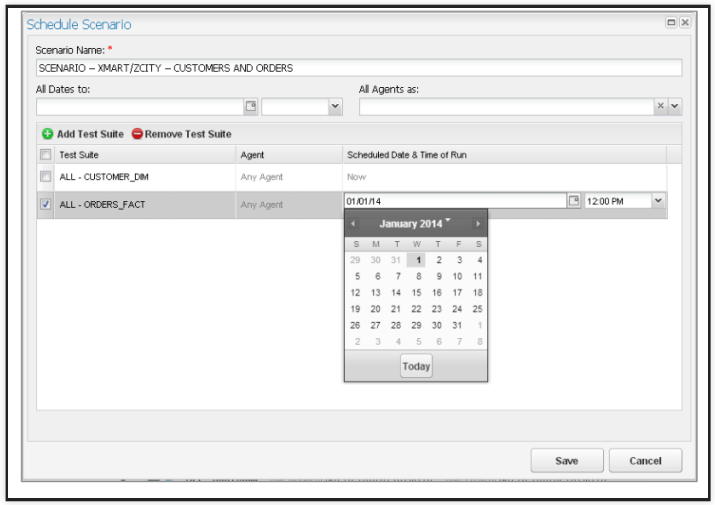

- Timing/Scheduling is a major cause of flakiness in Data Testing
- BUT - need to use inter-application/enterprise scheduler

QuerySurge Architecture

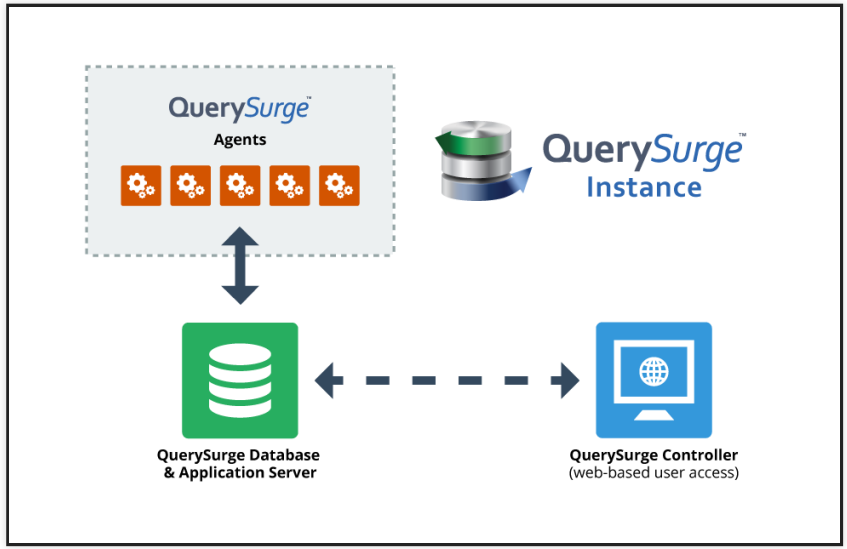


ready
MACHINERY MOVERS

Alternative Solution - Based on Unix Philosophy

- Write programs that do one thing and do it well.
- Write programs to work together.
- Write programs to handle text streams, because that is a universal interface.

Peter H. Salus, A Quarter-Century of Unix (1994)

```
fm@fm-VirtualBox:~/reveal.js$ diffst -t supplier_details
Setting up SUPPDB DB connection:...
Running SUPPDB SQL...
Setting up DW DB connection:...
Running DW SQL...
Fetching results...
+-----------------------------------------------+-----------------------------------------------+
|                                               >Supplier#000009976|ALGERIA|9975|Manufacturer#5    |
|Supplier#000009976|ALGERIA|62457|Manufacturer#5   <                                             |
|Supplier#000009976|ALGERIA|184939|Manufacturer#5  *Supplier#000009976|ALGERIA|184939|Manufactur     |
|Supplier#000009976|ALGERIA|197456|Manufacturer#9  *Supplier#000009976|ALGERIA|197456|Manufacturer#4  |
+-----------------------------------------------+-----------------------------------------------+
2019/08/05 16:38:38 Not OK - diffs found
fm@fm-VirtualBox:~/reveal.js$ diffst -t supplier_details -c
Setting up SUPPDB DB connection:...
Running SUPPDB SQL...
Setting up DW DB connection:...
Running DW SQL...
Fetching results...
+------------------------------+------------------------------+------------------------------+
|name                          |                              >Supplier#000009976            |
|country                       |                              >ALGERIA                       |
|partkey                       |                              >9975                          |
|mfgr                          |                              >Manufacturer#5                |
+------------------------------+------------------------------+------------------------------+
|name                          |Supplier#000009976            <                              |
|country                       |ALGERIA                       <                              |
|partkey                       |62457                         <                              |
|mfgr                          |Manufacturer#5                <                              |
+------------------------------+------------------------------+------------------------------+
|name                          |Supplier#000009976            |Supplier#000009976            |
|country                       |ALGERIA                       |ALGERIA                       |
|partkey                       |184939                        |184939                        |
|mfgr                          |Manufacturer#5                *Manufactur                    |
+------------------------------+------------------------------+------------------------------+
|name                          |Supplier#000009976            |Supplier#000009976            |
|country                       |ALGERIA                       |ALGERIA                       |
|partkey                       |197456                        |197456                        |
|mfgr                          |Manufacturer#9                *Manufacturer#4                |
+------------------------------+------------------------------+------------------------------+
2019/08/05 16:38:41 Not OK - diffs found
```

Limitations

- SQL-accessibility
- Limited by DB Drivers

- **Apache Ignite/GridGain**: https://github.com/amsokol/ignite-go-client
- **Apache Impala**: https://github.com/bippio/go-impala
- **Apache Avatica/Phoenix**: https://github.com/apache/calcite-avatica-go
- **AWS Athena**: https://github.com/segmentio/go-athena
- **ClickHouse** (uses native TCP interface): https://github.com/kshvakov/clickhouse
- **ClickHouse** (uses HTTP API): https://github.com/mailru/go-clickhouse
- **CockroachDB**: Use any PostgreSQL driver
- **Couchbase N1QL**: https://github.com/couchbase/go_n1ql
- **DB2 LUW** and **DB2/Z with DB2-Connect**: https://bitbucket.org/phiggins/db2cli (Last updated 2015-08)
- **DB2 LUW** (uses cgo): https://github.com/asifjalil/cli
- **DB2 LUW, z/OS, iSeries and Informix**: https://github.com/ibmdb/go_ibm_db
- **Firebird SQL**: https://github.com/nakagami/firebirdsql
- **MS ADODB**: https://github.com/mattn/go-adodb
- **MS SQL Server** (pure go): https://github.com/denisenkom/go-mssqldb
- **MS SQL Server** (uses cgo): https://github.com/minus5/gofreetds
- **MySQL**: https://github.com/ziutek/mymysql `[*]`
- **MySQL**: https://github.com/go-sql-driver/mysql/ `[*]`
- **ODBC**: https://bitbucket.org/miquella/mgodbc (Last updated 2016-02)
- **ODBC**: https://github.com/alexbrainman/odbc
- **Oracle**: https://github.com/mattn/go-oci8
- **Oracle**: https://gopkg.in/rana/ora.v4
- **Oracle**: https://gopkg.in/goracle.v2

Source and Target database options - 1

- **QL**: http://godoc.org/github.com/cznic/ql/driver
- **Postgres** (pure Go): https://github.com/lib/pq `[*]`
- **Postgres** (uses cgo): https://github.com/jbarham/gopgsqldriver
- **Postgres** (pure Go): https://github.com/jackc/pgx `[**]`
- **Presto**: https://github.com/prestodb/presto-go-client
- **SAP HANA** (uses cgo): https://help.sap.com/viewer/0eec0d68141541d1b07893a39944924e/2.0.03/en-US/0ffbe86c9d9f44338441829c6bee15e6.html
- **SAP HANA** (pure go): https://github.com/SAP/go-hdb
- **SAP ASE** (uses cgo): https://github.com/SAP/go-ase - package cgo (pure go package planned)
- **Snowflake** (pure Go): https://github.com/snowflakedb/gosnowflake
- **SQLite** (uses cgo): https://github.com/mattn/go-sqlite3 `[*]`
- **SQLite** (uses cgo): https://github.com/gwenn/gosqlite - Supports SQLite dynamic data typing
- **SQLite** (uses cgo): https://github.com/mxk/go-sqlite
- **SQLite**: (uses cgo): https://github.com/rsc/sqlite
- **SQL over REST**: https://github.com/adaptant-labs/go-sql-rest-driver
- **Sybase SQL Anywhere**: https://github.com/a-palchikov/sqlago
- **Sybase ASE** (pure go): https://github.com/thda/tds
- **Vitess**: https://godoc.org/vitess.io/vitess/go/vt/vitessdriver
- **YQL (Yahoo! Query Language)**: https://github.com/mattn/go-yql
- **Apache Hive**: https://github.com/sql-machine-learning/gohive
- **MaxCompute**: https://github.com/sql-machine-learning/gomaxcompute

Source and Target database options - 2

diffst - About

- Developed in Go (golang)
- Go programs compile into statically linked executables for Windows/Linux/Mac/...
- ⇒ No runtime required

Advantages of bespoke tool

- Extends the reach of manual testing
- Bug lifecycle simpler, faster, better
- CLI, numeric exit status, text input and output supports:
  - Job scheduling
  - Interface with test management tool via REST/SOAP API
  - Customizable test suites with scripting
  - Deployability - tester laptop or server
  - Test case management and maintainability thru version control

Questions/Comments?